Personalizing POI Descriptions with LLMs: A Study of User Perceptions and Informational Value

Anonymous Author(s)

Abstract

Most Points of Interest (POI) descriptions remain generic and static, presenting identical content to all users regardless of their interests or travel preferences. This limits perceived relevance and may reduce engagement with travel platforms. This study examines whether large language models (LLMs) can personalize POI descriptions to enhance perceived information quality. A within-subjects study was conducted with 51 participants who evaluated both original and LLM-adapted descriptions across 10 POIs spanning hotels, restaurants, museums, gardens, and landmarks. Participants rated descriptions on trustworthiness and clarity. Results indicate that AI-personalized descriptions significantly improved perceived clarity ($p = 0.0218$) and showed marginal improvements in trustworthiness ($p = 0.050$) compared to original descriptions. AI adaptations shifted user responses from neutral assessments toward stronger positive endorsements while maintaining engagement levels similar to original content (43.3% AI vs. 43.1% original; 13.5% equal). Participants also reported high comfort with AI-generated travel content (mean = 3.94, $n = 17$). These findings suggest that LLM-based personalization provides a scalable approach to improve tourism content relevance without compromising credibility.

CCS Concepts

• **Human-centered computing** → **User centered design**; *User studies.*

Keywords

poi adaptation, LLM-generated content, ai-driven descriptions, content personalization

ACM Reference Format:

Anonymous Author(s). 2018. Personalizing POI Descriptions with LLMs: A Study of User Perceptions and Informational Value. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX).* ACM, New York, NY, USA, 6 pages. https://doi.org/XXXXXXX.XXXXXXX

1 Introduction

Travel planning increasingly happens through digital platforms that help users evaluate accommodations, attractions, and other Points of Interest (POIs). Recommender systems have therefore become central in tourism applications, where personalization is used to cope with information overload and heterogeneous user needs [1, 8]. In the POI domain specifically, prior work has largely emphasized predicting which places users may prefer (e.g., next-POI or ranked POI recommendation) and how to incorporate richer signals such as context and heterogeneous data [6, 12]. Recent surveys further argue that large language models (LLMs) are becoming an enabling component across the recommendation pipeline, including user interaction and explanation generation [10, 11, 15].

However, while ranking and recommendation determine *which* POIs are shown, users still rely on POI *descriptions* to decide *why* a place is relevant and whether it is credible and worth visiting. In many systems, these descriptions remain static and generic, presenting the same framing to all readers despite differences in interests, travel styles, and constraints. This "one-size-fits-all" presentation can under-emphasize what matters most to a given user (e.g., historical context versus family suitability), potentially limiting perceived clarity and usefulness. Generative AI offers a way to adapt the *emphasis* of descriptions at scale, complementing POI recommendation and supporting more user-centered travel decision-making [5, 14].

At the same time, deploying LLM-generated travel content raises practical concerns around information quality and trust, especially in domains where credibility is critical. Tourism-focused GenAI research has begun to examine opportunities and risks (e.g., reliability, domain adaptation, and user acceptance), and approaches such as retrieval augmentation have been proposed to mitigate factual errors [4, 5, 7]. Motivated by this gap, this study examines whether LLM-based personalization of POI descriptions can improve perceived information quality without harming trust or engagement. A prototype was developed that presents users with original (manual) descriptions and LLM-adapted variants conditioned on stated preferences, and both versions were evaluated in a within-subjects study across multiple POIs. The contributions of this study are (1) an empirical comparison of manual versus LLM-adapted POI descriptions using significance, trustworthiness, and clarity as evaluation dimensions, and (2) evidence about user acceptance of adaptive description personalization in a tourism context.

2 Related work

Prior work on tourism and location-based personalization has largely focused on point-of-interest (POI) recommendation, where systems predict or rank places that match a user profile and context. Surveys of travel and POI recommenders highlight the importance of heterogeneous signals such as mobility patterns, time, geography, and user attributes, as well as evaluation challenges that arise from contextual variability and data sparsity [8, 12]. Recent work in the UMAP community also emphasizes that user preferences can differ before and after a visit, which complicates modeling and suggests that preference signals and content expectations are not static [3].

With the rise of large language models (LLMs), research has expanded from representation learning and ranking to interactive

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06
https://doi.org/XXXXXXX.XXXXXXX

and generative paradigms for personalization. In the POI domain, recent overviews and systems propose LLM-based pipelines that augment data, unify encoding, and support conversational interaction [10, 11]. Empirical studies and methods also explore how LLMs can be used for next-POI recommendation and for generating user profiles that act as conditioning context [6, 13]. In recommender systems more broadly, LLMs are increasingly studied as general-purpose personalization components that can support explanation, dialogue, and reasoning, but they also raise concerns about faithfulness and evaluation, which relate directly to perceived trust and clarity [2, 9, 15].

In parallel, tourism research has begun to evaluate generative AI for content creation and trip planning, including marketing materials and domain-specific models [4, 5, 14]. These works motivate the use of generative models in tourism, but they often study destination promotion or end-to-end planning rather than the personalization of the micro-content that users read when deciding whether a POI is relevant. This gap is addressed by adapting existing POI descriptions to a user-provided interest profile while keeping core facts stable, and by measuring user perceptions of significance, trustworthiness, and clarity in a controlled within-subjects study.

3 Methodology

This section outlines the methodology employed to evaluate the effectiveness of LLM-generated personalized POI descriptions compared to original manual descriptions. The study design, description personalization process, evaluation interface, metrics, procedure, and participant recruitment are described below.

3.1 Study Design

A within-subjects experimental design was employed in which each participant evaluated both original and AI-personalized descriptions for multiple POIs. This design enabled direct comparison between description types while controlling for individual differences in preferences and expectations. The study presented participants with 10 POIs spanning five categories: hotels, restaurants, museums, gardens, and landmarks. Each category included two POIs to ensure diverse representation of travel-related content. POIs were selected to include both internationally recognized locations and lesser-known destinations, providing variety in familiarity levels.

3.2 Description Personalization

For each POI, two description versions were prepared. The original description was sourced from official tourism websites and travel platforms, representing standard, non-personalized content. The AI-adapted version was generated using GPT-4o[1]. The model received a system prompt instructing it to act as an expert travel writer, emphasizing unique aspects, cultural significance, and visitor experience while adapting content to the user's situation. The user prompt included ten personalization dimensions: age, gender, marital status, whether the user has children, interests, travel experience level, education, profession, hobbies, and preferred travel style. The original POI title and description were provided, and the model was instructed to maintain factual accuracy while emphasizing aspects most relevant to the user's profile. For example, a museum description for a participant interested in history would emphasize historical significance and artifacts, while the same museum described for a family traveler would highlight interactive exhibits and educational programs.

[1]https://platform.openai.com/docs/models/gpt-4o

3.3 Evaluation Interface

Participants viewed descriptions through a web-based interface that displayed two versions side by side, labeled as Version A and Version B (see Figure 1). The interface did not indicate which version was AI-generated. Each POI was accompanied by a representative image to provide visual context. The presentation order of AI and manual descriptions was randomized across POIs to prevent order bias. Participants could not proceed to the next POI without completing evaluations for the current one, ensuring complete data collection.

3.4 Evaluation Metrics

Participants evaluated each description on three core dimensions using Likert scales:

Significance: "Does the description effectively communicate the significance and offerings of the place?" rated on a 5-point scale from Strongly Disagree to Strongly Agree. This metric assessed whether the description conveyed the value and appeal of the POI.

Trustworthiness: "How trustworthy does this description appear to be?" rated on a 5-point scale from Not at all to Extremely. This metric measured perceived credibility and reliability of the information.

Clarity: "How clear and complete is the information provided?" rated on a 5-point scale from Very Unclear to Very Clear. This metric evaluated information quality and comprehensibility.

After rating both descriptions individually, participants completed comparative evaluations selecting which description was more engaging, more informative, more appealing, and more likely to encourage them to visit. Participants also reported their familiarity with each POI using three categories: visited in person, heard about it, or never heard of it.

3.5 Procedure

The study was conducted in three phases using a secure web platform. First, participants provided informed consent and completed a survey regarding their demographics, travel styles, and interests. This data was used to personalize the AI-generated descriptions. Second, participants evaluated 10 POIs sequentially by reading both description versions and rating them on core metrics and familiarity. This evaluation phase lasted approximately 20 to 30 minutes. Finally, participants provided overall feedback on their experience and their comfort level with AI-generated travel content. All data were collected anonymously to ensure privacy. The study was conducted in accordance with institutional ethical guidelines, with participants providing informed consent prior to participation.

3.6 Participants

A total of 51 participants were recruited through personal networks via email and direct messaging. Participants were required to be at least 18 years old and possess basic English proficiency to ensure

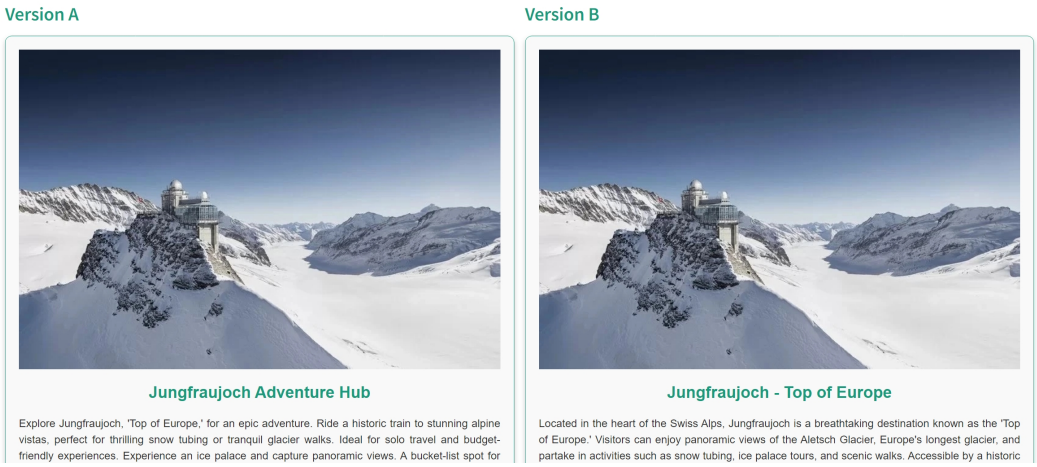

Figure 1: Evaluation interface showing two description versions (Version A and Version B) presented side by side for participant comparison

effective engagement with POI descriptions. No specific travel experience was mandated, as the study aimed to assess AI-generated descriptions across diverse user segments.

The sample demographics were as follows. Ages ranged from 22 to 49 years (M = 30.94, SD = 6.74). Gender distribution included 34 males (67%) and 17 females (33%). Nationalities were predominantly Indian (n = 28) and German (n = 21), with Chinese (n = 1) and Iranian (n = 1) representation. Marital status comprised 25 married, 22 single, 3 in relationships, and 1 divorced. Among the 44 participants who reported their profession, backgrounds centered on students (n = 22) and engineers (n = 16), with smaller numbers of academics (n = 3), healthcare professionals (n = 1), government workers (n = 1), and self-employed individuals (n = 1).

Travel experience levels showed intermediate (n = 21) and experienced (n = 19) travelers forming the majority, with 9 beginners and 2 experts. This distribution provided a range of perspectives on travel content while maintaining concentration in engaged, technology-familiar demographics.

4 Results

This section presents the findings from the user study evaluating LLM-adapted POI descriptions against original manual descriptions. Participant responses are analyzed across three core metrics: significance, trustworthiness, and clarity, followed by comparative preference evaluations and overall feedback on the adaptive personalization concept. The study collected 510 paired observations (51 participants × 10 POIs) for each metric, enabling robust statistical analysis of user perceptions. Table 1 provides an overview of the rating distributions across all three evaluation dimensions.

4.1 Perceived Significance

Participants rated both description versions on how effectively they communicated the significance and offerings of each POI using a 5-point Likert scale ranging from Strongly Disagree to Strongly Agree. Figure 2 illustrates the distribution of responses for manual and AI-generated descriptions.

Metric	Level (1–5)	Manual	AI
Significance	1 (Lowest)	1.4	1.6
	2	4.1	4.7
	3	27.1	23.1
	4	55.1	54.1
	5 (Highest)	12.3	16.5
Trust	1 (Lowest)	1.6	0.8
	2	8.4	6.5
	3	33.3	27.4
	4	46.1	51.2
	5 (Highest)	10.6	14.1
Clarity	1 (Lowest)	0.6	1.0
	2	6.5	4.7
	3	28.8	21.4
	4	51.0	55.7
	5 (Highest)	13.1	17.2

Table 1: Distribution of User Ratings % for Manual vs. AI Descriptions ($N = 510$)

The results show that both manual and AI descriptions received predominantly positive ratings. In both conditions, the largest share of responses fell in level 4 (Agree; Manual: 55.1%, AI: 54.1%), with additional responses at the highest level 5 (Strongly Agree; Manual: 12.3%, AI: 16.5%). AI-generated descriptions exhibited a shift toward stronger endorsement, reflected in a 4.2 percentage-point increase in top-box ratings (level 5) and a corresponding reduction in neutral responses (level 3; Manual: 27.1%, AI: 23.1%). A chi-square test on the underlying $N = 510$ observations revealed no statistically significant difference between distributions ($\chi^2(4) = 4.87$, $p = 0.3005$), indicating that AI adaptations maintained comparable perceived significance while trending toward higher extreme-positive ratings.

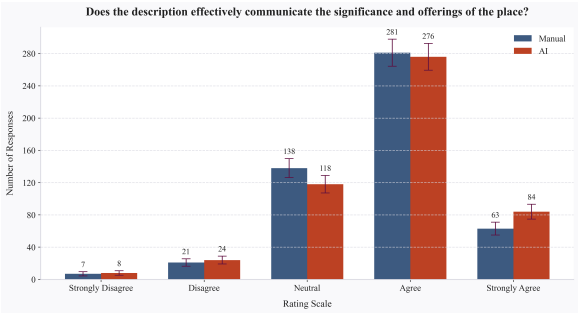

Figure 2: Comparison of perceived significance between manual and AI-generated descriptions. Error bars represent Poisson-style standard error estimates computed on raw counts ($\sqrt{n}$).

4.2 Trustworthiness

Trustworthiness was assessed on a 5-point scale ranging from Not at all to Extremely. Figure 3 presents the comparative distribution of trust ratings.

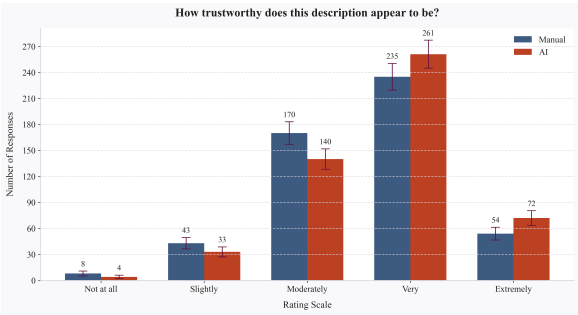

Figure 3: Comparison of perceived trustworthiness between manual and AI-generated descriptions. Error bars represent Poisson-style standard error estimates computed on raw counts ($\sqrt{n}$).

AI-generated descriptions received higher trust ratings overall than manual descriptions. At the top end of the scale, the share of "Extremely" trustworthy responses (level 5) increased from 10.6% (manual) to 14.1% (AI), while "Very" trustworthy responses (level 4) rose from 46.1% to 51.2%. At the lower end, "Not at all" trustworthy ratings (level 1) decreased from 1.6% to 0.8%, and level 2 ratings also declined (8.4% to 6.5%). This pattern reflects a shift away from moderate trust (level 3; 33.3% to 27.4%) toward higher trust categories. A chi-square test showed a borderline association between condition and trustworthiness ratings ($\chi^2(4) = 9.49$, $p = .050$).

4.3 Clarity and Completeness

Participants evaluated how clear and complete the information was using a 5-point scale from Very Unclear to Very Clear. Figure 4 shows the distribution of clarity ratings.

AI-generated descriptions demonstrated superior clarity and completeness ratings compared to manual versions. The proportion

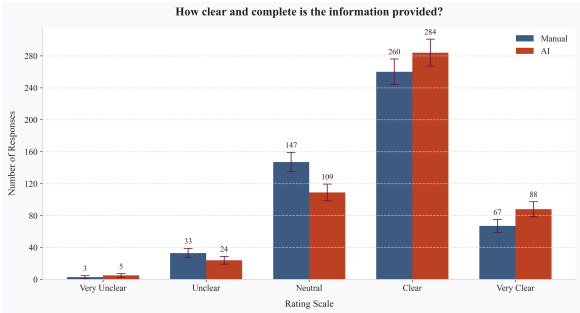

Figure 4: Comparison of perceived clarity and completeness between manual and AI-generated descriptions. Error bars represent Poisson-style standard error estimates computed on raw counts ($\sqrt{n}$).

of responses at level 4 ("Clear") increased from 51.0% (manual) to 55.7% (AI), and level 5 ("Very Clear") increased from 13.1% to 17.2%. This improvement coincided with a substantial reduction in neutral responses (level 3; 28.8% to 21.4%) and a decrease at level 2 (6.5% to 4.7%), indicating that AI adaptations helped participants form more definitive positive judgments about information quality. A chi-square test confirmed statistical significance ($\chi^2(4) = 11.47$, $p = 0.0218$), providing evidence that personalized AI descriptions significantly enhanced perceived clarity and completeness.

4.4 Comparative Preferences

Beyond individual ratings, participants indicated which description version they found more engaging, relevant, and appealing. For the engaging preference, responses were nearly evenly split between AI-generated (221; 43.3%) and manual (220; 43.1%) descriptions, with 69 selections (13.5%) indicating both versions were equally engaging. This near-parity suggests that AI personalization matched the engagement level of carefully crafted manual descriptions while adding the benefit of customization.

4.5 Overall Study Feedback

Following the POI evaluations, participants were invited to provide optional overall feedback. A subset of 17 participants completed this section, rating three dimensions using a 5-point scale (1 = Poor, 5 = Excellent). The overall experience with POI descriptions received a mean rating of 3.88 (SD = 0.70, Median = 4.0), indicating generally positive reception. Participants expressed moderate to strong support for the concept of automatically adapting POI descriptions based on user interests, with a mean rating of 3.75 (SD = 0.93, Median = 4.0). Comfort with reading AI-generated descriptions when planning visits to new places received the highest mean rating of 3.94 (SD = 0.83, Median = 4.0), suggesting participants were receptive to AI-generated travel content when it demonstrated quality and relevance. The relatively low standard deviations (all below 1.0) indicate consensus among participants without extreme polarization of opinions.

5 Discussion

This study examined whether LLM-based personalization of POI descriptions changes how people judge the quality of travel information. Overall, participants rated the AI-adapted descriptions as clearer than the manual baseline descriptions. Ratings of trustworthiness showed a borderline improvement ($p = 0.050$), suggesting that AI adaptations at minimum maintained credibility while trending toward higher trust ratings. Participants' choices of which description they would prefer to read were also broadly stable. Taken together, these results suggest that personalization can increase perceived clarity without reducing willingness to engage with the content.

Across measures, AI-adapted descriptions more often received positive ratings rather than neutral ratings. A likely reason is increased *perceived relevance*. When a description highlights details that match a reader's interests, the same facts may feel more useful and easier to understand. This interpretation fits prior work on tailored communication, where framing and emphasis influence how information is processed even when the underlying content is unchanged. At the same time, the mostly stable trust ratings suggest that rewriting alone may not strongly change credibility judgments in a low-stakes reading task.

From a practical perspective, these findings support the use of LLMs to adjust emphasis and framing in POI micro-content using lightweight preference signals (e.g., stated interests). This approach can help scale personalization without requiring extensive manual editing. However, trust is sensitive to context and presentation. Platforms should therefore treat transparency and disclosure as design choices and aim to set clear expectations without unnecessarily signaling low quality.

5.1 Limitations

Several limitations should be acknowledged. First, participants evaluated descriptions in a controlled study setting rather than during real-world trip planning, which may limit ecological validity. Second, the study focused on perceived outcomes (e.g., clarity and trustworthiness) and did not measure downstream behavior such as click-through, saving, itinerary inclusion, or actual visits. Third, the within-subjects design supports direct comparison but may increase attention to differences between versions and introduce demand characteristics. Fourth, the study did not test potential downsides of personalization, such as reduced serendipity or overfitting to stated interests. Finally, the sample composition may limit generalizability to other age groups, levels of technical familiarity, and cultural contexts.

6 Future Work

Future research directions include expanding the study to include more diverse participants, such as older adults and people from different cultural backgrounds. Such expansion would help clarify how age and culture affect user acceptance of AI-generated travel content. Beyond short-term surveys, longitudinal field studies with real travel platforms should be conducted to track actual user behavior, including booking rates and long-term trust. Additionally, more detailed personalization factors such as trip purpose, travel companions, and budget constraints could be tested to examine their effects on user satisfaction.

A second research direction involves the transparency and integration of AI-generated descriptions. Investigation is needed into how explaining why certain information was highlighted can improve user trust. It is also important to study the potential for filter bubbles where over-personalization might limit new and surprising discoveries. Methods that balance personalized relevance with novelty and diversity should be developed. Finally, the integration of adaptive descriptions into existing recommendation systems warrants exploration, particularly to support users from the start of their search until booking completion.

7 Conclusion

This study shows that large language models can personalize point-of-interest descriptions in ways that improve how users read and judge travel information. Compared to non-personalized text, the AI-adapted descriptions were generally perceived as clearer, and they also tended to be seen as more trustworthy. Participants also reported feeling comfortable engaging with AI-generated content in this setting.

These findings suggest that personalization can help people make more confident decisions by foregrounding details that match their interests while preserving the core facts of a place. For travel platforms, LLM-based rewriting offers a practical path to scaling tailored micro-content without relying on extensive manual editorial effort. More broadly, the study highlights that how a place is described can be as influential as which place is recommended for creating a relevant and engaging user experience.

References

[1] Gediminas Adomavicius and Alexander Tuzhilin. 2005. Toward the Next Generation of Recommender Systems: A Survey of the State-of-the-Art and Possible Extensions. *IEEE Transactions on Knowledge and Data Engineering* 17, 6 (2005), 734–749. doi:10.1109/TKDE.2005.99

[2] Yue Feng, Shuchang Liu, Zhenghai Xue, Qingpeng Cai, Lantao Hu, Peng Jiang, Kun Gai, and Fei Sun. 2023. A Large Language Model Enhanced Conversational Recommender System. doi:10.48550/arXiv.2308.06212 arXiv:2308.06212

[3] Katharina Hofschen, David Massimo, and Francesco Ricci. 2023. Expected and Experienced Utility of Points of Interest in Tourism Recommender Systems. In *Adjunct Proceedings of the 31st ACM Conference on User Modeling, Adaptation and Personalization* (New York, NY, USA, 2023-06-16) *(UMAP '23 Adjunct)*. Association for Computing Machinery, 50–55. doi:10.1145/3563359.3597405

[4] Cathy H. C. Hsu, Guoxiong Tan, and Bela Stantic. 2024. A fine-tuned tourism-specific generative AI concept. *Annals of Tourism Research* 104 (2024), 103723. doi:10.1016/j.annals.2023.103723

[5] Hengyun Li, Jingbo Xi, Cathy H. C. Hsu, Bruce X. B. Yu, and Xiang (Kevin) Zheng. 2025. Generative artificial intelligence in tourism management: An integrative review and roadmap for future research. *Tourism Management* 110 (2025), 105179. doi:10.1016/j.tourman.2025.105179

[6] Peibo Li, Maarten de Rijke, Hao Xue, Shuang Ao, Yang Song, and Flora D. Salim. 2024. Large Language Models for Next Point-of-Interest Recommendation. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '24)*. Association for Computing Machinery, 1463–1472. doi:10.1145/3626772.3657840

[7] Jinhu Qi, Shuai Yan, Yibo Zhang, Wentao Zhang, Rong Jin, Yuwei Hu, and Ke Wang. 2025. RAG-Optimized Tibetan Tourism LLMs: Enhancing Accuracy and Personalization. doi:10.48550/arXiv.2408.12003 arXiv:2408.12003 [cs]

[8] Shini Renjith, A. Sreekumar, and M. Jathavedan. 2020. An extensive study on the evolution of context-aware personalized travel recommender systems. *Information Processing & Management* 57, 1 (2020), 102078. doi:10.1016/j.ipm.2019.102078

[9] Alicia Tsai, Adam Kraft, Long Jin, Chenwei Cai, Anahita Hosseini, Taibai Xu, Zemin Zhang, Lichan Hong, Ed Chi, and Xinyang Yi. 2024. Leveraging LLM Reasoning Enhances Personalized Recommender Systems. In *Findings of the Association for Computational Linguistics: ACL 2024*. Association for Computational

Linguistics, 13176–13188. doi:10.18653/v1/2024.findings-acl.780

[10] Tianxing Wang and Can Wang. 2024. Embracing LLMs for Point-of-Interest Recommendations. *IEEE Intelligent Systems* 39, 1 (2024), 56–59. doi:10.1109/MIS.2023.3343489

[11] Tianxing Wang and Can Wang. 2025. Point-of-Interest Recommendations and Large Language Models: A Powerful Combination. *IEEE Intelligent Systems* 40, 3 (2025), 5–9. doi:10.1109/MIS.2025.3561473

[12] Zehui Wang, Wolfram Höpken, and Dietmar Jannach. 2025. A survey on point-of-interest recommendations leveraging heterogeneous data. *Information Technology & Tourism* 27, 1 (2025), 29–73. doi:10.1007/s40558-024-00301-3

[13] Wilson Wongso, Hao Xue, and Flora Salim. 2025. GenUP: Generative User Profilers as In-Context Learners for Next POI Recommender Systems. In *Proceedings of the 33rd ACM International Conference on Advances in Geographic Information Systems* (2025-12-12) *(SIGSPATIAL '25)*. Association for Computing Machinery, 436–439. doi:10.1145/3748636.3762754

[14] Yaozhi Zhang and Nina Katrine Prebensen. 2024. Co-creating with ChatGPT for tourism marketing materials. *Annals of Tourism Research Empirical Insights* 5, 1 (2024), 100124. doi:10.1016/j.annale.2024.100124

[15] Zihuai Zhao, Wenqi Fan, Jiatong Li, Yunqing Liu, Xiaowei Mei, Yiqi Wang, Zhen Wen, Fei Wang, Xiangyu Zhao, Jiliang Tang, and Qing Li. 2024. Recommender Systems in the Era of Large Language Models (LLMs). *IEEE Transactions on Knowledge and Data Engineering* 36, 11 (2024), 6889–6907. doi:10.1109/TKDE.2024.3392335

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009