Personalizing POI Descriptions with LLMs: A Study of User Perceptions and Informational Value

Anonymous Author(s)

Abstract

Personalized content is essential for engaging users in travel planning, yet most Points of Interest (POI) descriptions remain generic and static. This study examines whether large language models (LLMs) can effectively personalize POI descriptions to enhance user engagement and perceived information quality. A within-subjects study was conducted with 51 participants who evaluated both original and LLM-adapted descriptions across 10 diverse POIs, including hotels, restaurants, museums, and gardens. Participants rated descriptions for significance, trustworthiness, and clarity. Results indicate that AI-personalized descriptions significantly improved perceived clarity ($p = 0.0218$) and showed marginal improvements in trustworthiness ($p = 0.050$) compared to manual descriptions. AI adaptations shifted user responses from neutral assessments toward stronger positive endorsements while maintaining engagement levels similar to carefully crafted manual content. Participants also reported high comfort with AI-generated travel content (mean = 3.94 on a 5-point scale). These findings suggest that LLM-based personalization provides a scalable approach to improve tourism content relevance and user engagement without compromising credibility or information quality.

CCS Concepts

• **Human-centered computing** → **User centered design**; *User studies.*

Keywords

poi adaptation, LLM-generated content, ai-driven descriptions, content personalization

ACM Reference Format:

Anonymous Author(s). 2018. Personalizing POI Descriptions with LLMs: A Study of User Perceptions and Informational Value. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX).* ACM, New York, NY, USA, **??** pages. https://doi.org/XXXXXXX.XXXXXXX

1 Introduction

Personalization has become a key factor in improving user engagement in various digital experiences, including travel and tourism. As travellers seek more relevant and tailored recommendations, AI advancements, particularly large language models (LLM), presents new opportunities to adapt content to align with individual preferences dynamically. One such application is personalized descriptions of Points of Interest (POIs), such as hotels, restaurants, museums, and gardens. Although traditional POI descriptions are often static and generalized, AI-driven personalization can refine these descriptions to match user-specific interests, past experiences, and contextual relevance, enhancing both engagement and comprehension.

Much prior research has focused on POI recommendation and ranking systems, where AI models suggest locations based on user preferences, behaviors, and contextual factors. However, less attention has been given to how the content of each POI itself can be adapted to better engage users. When individuals access POI descriptions, the same information is typically presented to all users, despite variations in their backgrounds, interests, and motivations to visit a place. For example, a museum description could emphasize historical significance for a history enthusiast while focusing on interactive exhibits for a family traveler. Generative AI and LLMs provide a unique opportunity to dynamically tailor content without altering its factual meaning, ensuring that the description remains relevant and appealing to different audiences.

To explore this potential, a prototype system was developed that allows users to enter their preferences and interests and then dynamically presents two versions of each POI description: the original and an LLM-adapted version. Participants were asked to evaluate the descriptions based on alignment with their preferences, perceived appeal, and familiarity with the location. The key objective is to assess whether AI-generated adaptations are acceptable and meaningful, effectively convey the original value of POI, and improve user engagement.

This study contributes to the growing field of AI-powered content personalization in tourism by investigating the role of LLMs in dynamic adaptation of textual content. By analyzing participant feedback, this study aims to provide insight into how AI-generated descriptions can improve engagement while preserving accuracy and authenticity.

2 Related work

Prior work on tourism and location-based personalization has largely focused on point-of-interest (POI) recommendation, where systems predict or rank places that match a user profile and context. Surveys of travel and POI recommenders highlight the importance of heterogeneous signals such as mobility patterns, time, geography, and user attributes, as well as evaluation challenges that arise from contextual variability and data sparsity [**? ?**]. Recent work in the UMAP community also emphasizes that user preferences can differ before and after a visit, which complicates modeling and suggests that preference signals and content expectations are not static [**?**].

With the rise of large language models (LLMs), research has expanded from representation learning and ranking to interactive and

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06
https://doi.org/XXXXXXX.XXXXXXX

generative paradigms for personalization. In the POI domain, recent overviews and systems propose LLM-based pipelines that augment data, unify encoding, and support conversational interaction [? ?]. Empirical studies and methods also explore how LLMs can be used for next-POI recommendation and for generating user profiles that act as conditioning context [? ?]. In recommender systems more broadly, LLMs are increasingly studied as general-purpose personalization components that can support explanation, dialogue, and reasoning, but they also raise concerns about faithfulness and evaluation, which relate directly to perceived trust and clarity [? ? ?].

In parallel, tourism research has begun to evaluate generative AI for content creation and trip planning, including marketing materials and domain-specific models [? ? ?]. These works motivate the use of generative models in tourism, but they often study destination promotion or end-to-end planning rather than the personalization of the micro-content that users read when deciding whether a POI is relevant. This gap is addressed by adapting existing POI descriptions to a user-provided interest profile while keeping core facts stable, and by measuring user perceptions of significance, trustworthiness, and clarity in a controlled within-subjects study.

3 Methodology

To evaluate the effectiveness of LLM-based personalization of POI descriptions, a controlled user study was conducted comparing original and AI-adapted descriptions. This section details the study design, POI selection, personalization approach, evaluation procedure, and participant characteristics.

3.1 Study Design

A within-subjects experimental design was employed in which each participant evaluated both original and AI-personalized descriptions for multiple POIs. This design enabled direct comparison between description types while controlling for individual differences in preferences and expectations. The study presented participants with 10 POIs spanning five categories: hotels, restaurants, museums, gardens, and landmarks. Each category included two POIs to ensure diverse representation of travel-related content. POIs were selected to include both internationally recognized locations and lesser-known destinations, providing variety in familiarity levels.

3.2 Description Personalization

For each POI, two description versions were prepared. The original description was sourced from official tourism websites and travel platforms, representing standard, non-personalized content. The AI-adapted version was generated using a large language model (GPT-4) that received the original description along with each participant's stated interests, hobbies, and travel preferences as input. The LLM was instructed to adapt the description to emphasize aspects most relevant to the user's profile while maintaining factual accuracy and the core information about the POI. For example, a museum description for a participant interested in history would emphasize historical significance and artifacts, while the same museum described for a family traveler would highlight interactive exhibits and educational programs.

3.3 Evaluation Interface

Participants viewed descriptions through a web-based interface that displayed two versions side by side, labeled as Version A and Version B (see Figure **??**). The interface did not indicate which version was AI-generated. Each POI was accompanied by a representative image to provide visual context. The presentation order of AI and manual descriptions was randomized across POIs to prevent order bias. Participants could not proceed to the next POI without completing evaluations for the current one, ensuring complete data collection.

3.4 Evaluation Metrics

Participants evaluated each description on three core dimensions using Likert scales:

Significance: "Does the description effectively communicate the significance and offerings of the place?" rated on a 5-point scale from Strongly Disagree to Strongly Agree. This metric assessed whether the description conveyed the value and appeal of the POI.

Trustworthiness: "How trustworthy does this description appear to be?" rated on a 5-point scale from Not at all to Extremely. This metric measured perceived credibility and reliability of the information.

Clarity: "How clear and complete is the information provided?" rated on a 5-point scale from Very Unclear to Very Clear. This metric evaluated information quality and comprehensibility.

After rating both descriptions individually, participants completed comparative evaluations selecting which description was more engaging, more informative, more appealing, and more likely to encourage them to visit. Participants also reported their familiarity with each POI using three categories: visited in person, heard about it, or never heard of it.

3.5 Procedure

The study was conducted in three phases using a secure web platform. First, participants provided informed consent and completed a survey regarding their demographics, travel styles, and interests. This data was used to personalize the AI-generated descriptions. Second, participants evaluated 10 POIs sequentially by reading both description versions and rating them on core metrics and familiarity. This evaluation phase lasted approximately 20 to 30 minutes. Finally, participants provided overall feedback on their experience and their comfort level with AI-generated travel content. All data were collected anonymously to ensure privacy and ethical compliance.

3.6 Participants

A total of 51 participants were recruited through personal networks via email and direct messaging. Participants were required to be at least 18 years old and possess basic English proficiency to ensure effective engagement with POI descriptions. No specific travel experience was mandated, as the study aimed to assess AI-generated descriptions across diverse user segments.

The sample demographics were as follows. Ages ranged from 22 to 49 years (M = 30.94, SD = 6.74). Gender distribution included 34 males (67%) and 17 females (33%). Nationalities were predominantly Indian (n = 28) and German (n = 21), with minimal Chinese and Iranian representation. Marital status comprised 25 married, 22

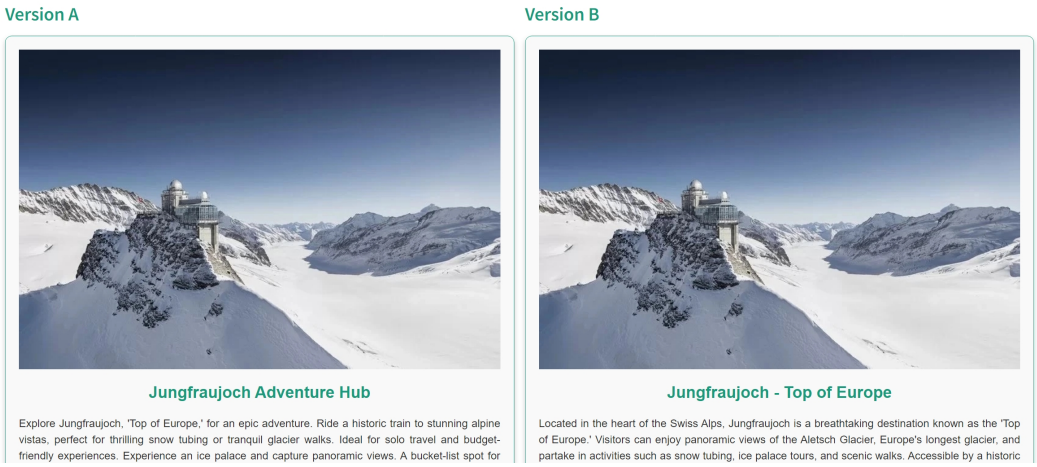

Figure 1: Evaluation interface showing two description versions (Version A and Version B) presented side by side for participant comparison

single, 3 in relationships, and 1 divorced. Professional backgrounds centered on students (n = 22) and engineers (n = 16), with smaller numbers of academics, healthcare professionals, and public service workers.

Travel experience levels showed intermediate (n = 21) and experienced (n = 19) travelers forming the majority, with 9 beginners and 2 experts. This distribution provided a range of perspectives on travel content while maintaining concentration in engaged, technology-familiar demographics.

4 Results

This section presents the findings from the user study evaluating LLM-adapted POI descriptions against original manual descriptions. Participant responses are analyzed across three core metrics: significance, trustworthiness, and clarity, followed by comparative preference evaluations and overall feedback on the adaptive personalization concept. The study collected 510 paired observations (51 participants × 10 POIs) for each metric, enabling robust statistical analysis of user perceptions. Table **??** provides an overview of the rating distributions across all three evaluation dimensions.

4.1 Perceived Significance

Participants rated both description versions on how effectively they communicated the significance and offerings of each POI using a 5-point Likert scale ranging from Strongly Disagree to Strongly Agree. Figure **??** illustrates the distribution of responses for manual and AI-generated descriptions.

The results show that both manual and AI descriptions received predominantly positive ratings, with the majority of responses falling in the Agree category (Manual: 281, AI: 276). However, AI-generated descriptions demonstrated a notable shift toward stronger agreement, with 84 Strongly Agree responses compared to 63 for manual descriptions. This increase was primarily driven by a reduction in Neutral responses (Manual: 138, AI: 118), suggesting that personalized descriptions moved participants from ambivalence toward stronger endorsement. A chi-square test revealed no statistically significant difference between the distributions ($\chi^2(4) = 4.87$, $p = 0.3005$), indicating that AI adaptations maintained equivalence with manual descriptions while showing a trend toward higher top-box scores.

Table 1: Summary of rating distributions for manual and AI-generated descriptions across three evaluation metrics (n=510 per condition)

Metric	Rating	Manual	AI
Significance	Strongly Disagree	7	8
	Disagree	21	24
	Neutral	138	118
	Agree	281	276
	Strongly Agree	63	84
Trustworthiness	Not at all	8	4
	Slightly	43	33
	Moderately	170	140
	Very	235	261
	Extremely	54	72
Clarity	Very Unclear	3	5
	Unclear	33	24
	Neutral	147	109
	Clear	260	284
	Very Clear	67	88

4.2 Trustworthiness

Trustworthiness was assessed on a 5-point scale ranging from Not at all to Extremely. Figure **??** presents the comparative distribution of trust ratings.

AI-generated descriptions received notably higher trust ratings than manual descriptions. The most striking difference appeared

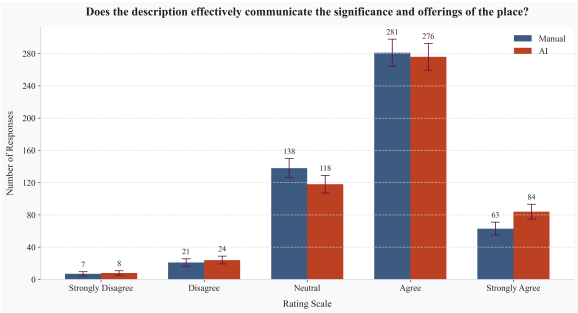

Figure 2: Comparison of perceived significance between manual and AI-generated descriptions. Error bars represent Poisson standard error estimates ($\sqrt{n}$).

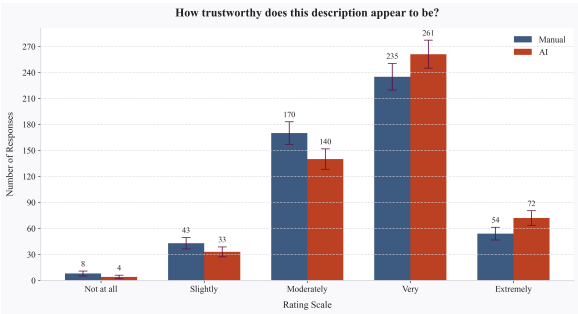

Figure 3: Comparison of perceived trustworthiness between manual and AI-generated descriptions. Error bars represent Poisson standard error estimates ($\sqrt{n}$).

at the highest end of the scale, with AI descriptions receiving 72 Extremely trustworthy ratings compared to 54 for manual descriptions. Additionally, AI descriptions garnered 261 Very trustworthy ratings versus 235 for manual versions. At the lower end of the scale, AI descriptions received only 4 Not at all ratings compared to 8 for manual descriptions, representing a 50% reduction in the most skeptical responses. The shift from Moderate trust toward Very and Extremely trust was particularly pronounced. A chi-square test indicated marginal statistical significance ($\chi^2(4) = 9.49$, $p = 0.050$), suggesting that AI personalization had a meaningful positive impact on perceived trustworthiness.

4.3 Clarity and Completeness

Participants evaluated how clear and complete the information was using a 5-point scale from Very Unclear to Very Clear. Figure **??** shows the distribution of clarity ratings.

AI-generated descriptions demonstrated superior clarity ratings compared to manual versions. Clear ratings increased from 260 for manual descriptions to 284 for AI descriptions, while Very Clear ratings rose from 67 to 88. This improvement was accompanied by a substantial reduction in Neutral responses (from 147 to 109), indicating that AI adaptations helped participants form more definitive judgments about information quality. The Unclear category also

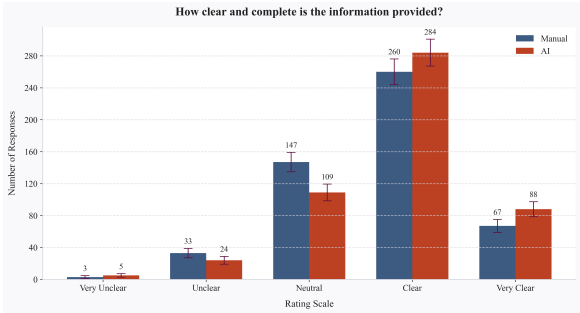

Figure 4: Comparison of perceived clarity and completeness between manual and AI-generated descriptions. Error bars represent Poisson standard error estimates ($\sqrt{n}$).

decreased from 33 to 24 responses. A chi-square test confirmed statistical significance ($\chi^2(4) = 11.47$, $p = 0.0218$), providing evidence that personalized AI descriptions significantly enhanced perceived clarity and completeness of information.

4.4 Comparative Preferences

Beyond individual ratings, participants indicated which description version they found more engaging, relevant, and appealing. For the engaging preference, responses were nearly evenly split between AI-generated (221 preferences) and manual (220 preferences), with 69 participants indicating both versions were equally engaging. This near-parity suggests that AI personalization successfully matched the engagement level of carefully crafted manual descriptions while adding the benefit of customization.

4.5 Overall Study Feedback

Following the POI evaluations, a subset of 17 participants provided overall ratings on three dimensions using a 5-point scale (1 = Poor, 5 = Excellent). The overall experience with POI descriptions received a mean rating of 3.88 (SD = 0.70, Median = 4.0), indicating generally positive reception. Participants expressed moderate to strong support for the concept of automatically adapting POI descriptions based on user interests, with a mean rating of 3.75 (SD = 0.93, Median = 4.0). Comfort with reading AI-generated descriptions when planning visits to new places received the highest mean rating of 3.94 (SD = 0.83, Median = 4.0), suggesting participants were receptive to AI-generated travel content when it demonstrated quality and relevance. The relatively low standard deviations (all below 1.0) indicate consensus among participants without extreme polarization of opinions.

5 Discussion

This study examined whether LLM-generated personalization of POI descriptions improves user engagement and perceived information quality without undermining trust. Overall, AI-adapted descriptions were rated as clearer than manual descriptions (significant improvement; $p = 0.0218$) and showed comparable trustworthiness with a marginal improvement in trust ratings ($p = 0.050$).

Engagement preferences were essentially unchanged across conditions (221 AI vs. 220 manual), suggesting that personalization can enhance perceived quality without reducing willingness to engage.

The pattern of ratings indicates that personalization primarily shifts evaluations away from neutrality toward more positive judgments across significance, trustworthiness, and clarity. Trust is a critical factor in travel decision-making, and AI-adapted descriptions received fewer "Not at all trustworthy" ratings and more "Extremely trustworthy" ratings than manual versions. A plausible mechanism is that aligning emphasis with stated interests increases perceived relevance while keeping the underlying facts stable; similarly, highlighting reader-relevant details may make the same information easier to process, yielding higher clarity without necessarily altering content.

Generalizability is limited by the sample, which was predominantly young and educated (mean age 30.94 years), with many students and engineers, and primarily Indian and German participants; responses may differ for older users, less tech-familiar groups, or other cultural contexts. For tourism platforms, the results suggest that LLMs can tailor emphasis in POI descriptions based on onboarding inputs or inferred preferences while maintaining perceived trust and improving clarity. Given the sensitivity of trust, any disclosure of AI involvement should be designed to set appropriate expectations without unnecessarily eroding credibility.

5.1 Limitations

Several limitations apply. Descriptions were evaluated in a controlled setting rather than during real trip planning. Behavioral outcomes were not measured beyond stated preferences. The within-subjects design may have increased attention to differences between versions. Potential harms of over-personalization (e.g., filter bubbles and reduced serendipitous discovery) were not evaluated.

6 Future Work

Future research will first expand the study to include more diverse participants, such as older adults and people from different cultural backgrounds. This will help us understand how age and culture affect how users accept AI-generated travel content. We also plan to move beyond short-term surveys by conducting longitudinal field studies with real travel platforms. These studies will track actual user behavior, including booking rates and long-term trust. Additionally, we will test more detailed personalization by adding factors like trip purpose, travel companions, and budget constraints to see how they affect user satisfaction.

A second research direction involves the transparency and integration of these AI descriptions. We will investigate how explaining why certain information was highlighted can improve user trust. It is also important to study the potential for filter bubbles where over-personalization might limit new and surprising discoveries. We plan to develop methods that balance personalized relevance with novelty and diversity. Finally, we will explore how these adaptive descriptions can be integrated into existing recommendation systems to support users from the start of their search until they complete a booking.

7 Conclusion

This study demonstrates that Large Language Models can effectively personalize POI descriptions to improve user engagement and information quality. A study of 51 participants and 510 comparisons across 10 POIs found that AI-generated adaptations significantly improved clarity ($p = 0.0218$) and showed marginal gains in trustworthiness ($p = 0.050$). The results suggest that personalization helps users form more confident assessments by highlighting details that match their interests, with participants reporting high comfort using AI-generated content (mean = 3.94/5). These findings show that travel platforms can use LLMs to scale personalized content without the high costs of manual writing. Ultimately, this research highlights that adapting how content is described is just as important as the recommendation itself for creating a relevant and engaging user experience.

References

[] Yue Feng, Shuchang Liu, Zhenghai Xue, Qingpeng Cai, Lantao Hu, Peng Jiang, Kun Gai, and Fei Sun. 2023. A Large Language Model Enhanced Conversational Recommender System. doi:10.48550/arXiv.2308.06212 arXiv:2308.06212

[] Katharina Hofschen, David Massimo, and Francesco Ricci. 2023. Expected and Experienced Utility of Points of Interest in Tourism Recommender Systems. In *Adjunct Proceedings of the 31st ACM Conference on User Modeling, Adaptation and Personalization* (New York, NY, USA, 2023-06-16) *(UMAP '23 Adjunct)*. Association for Computing Machinery, 50–55. doi:10.1145/3563359.3597405

[] Cathy H. C. Hsu, Guoxiong Tan, and Bela Stantic. 2024. A fine-tuned tourism-specific generative AI concept. *Annals of Tourism Research* 104 (2024), 103723. doi:10.1016/j.annals.2023.103723

[] Hengyun Li, Jingbo Xi, Cathy H. C. Hsu, Bruce X. B. Yu, and Xiang (Kevin) Zheng. 2025. Generative artificial intelligence in tourism management: An integrative review and roadmap for future research. *Tourism Management* 110 (2025), 105179. doi:10.1016/j.tourman.2025.105179

[] Peibo Li, Maarten de Rijke, Hao Xue, Shuang Ao, Yang Song, and Flora D. Salim. 2024. Large Language Models for Next Point-of-Interest Recommendation. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '24)*. Association for Computing Machinery, 1463–1472. doi:10.1145/3626772.3657840

[] Shini Renjith, A. Sreekumar, and M. Jathavedan. 2020. An extensive study on the evolution of context-aware personalized travel recommender systems. *Information Processing & Management* 57, 1 (2020), 102078. doi:10.1016/j.ipm.2019.102078

[] Alicia Tsai, Adam Kraft, Long Jin, Chenwei Cai, Anahita Hosseini, Taibai Xu, Zemin Zhang, Lichan Hong, Ed Chi, and Xinyang Yi. 2024. Leveraging LLM Reasoning Enhances Personalized Recommender Systems. In *Findings of the Association for Computational Linguistics: ACL 2024*. Association for Computational Linguistics, 13176–13188. doi:10.18653/v1/2024.findings-acl.780

[] Tianxing Wang and Can Wang. 2024. Embracing LLMs for Point-of-Interest Recommendations. *IEEE Intelligent Systems* 39, 1 (2024), 56–59. doi:10.1109/MIS.2023.3343489

[] Tianxing Wang and Can Wang. 2025. Point-of-Interest Recommendations and Large Language Models: A Powerful Combination. *IEEE Intelligent Systems* 40, 3 (2025), 5–9. doi:10.1109/MIS.2025.3561473

[] Zehui Wang, Wolfram Höpken, and Dietmar Jannach. 2025. A survey on point-of-interest recommendations leveraging heterogeneous data. *Information Technology & Tourism* 27, 1 (2025), 29–73. doi:10.1007/s40558-024-00301-3

[] Wilson Wongso, Hao Xue, and Flora Salim. 2025. GenUP: Generative User Profilers as In-Context Learners for Next POI Recommender Systems. In *Proceedings of the 33rd ACM International Conference on Advances in Geographic Information Systems* (2025-12-12) *(SIGSPATIAL '25)*. Association for Computing Machinery, 436–439. doi:10.1145/3748636.3762754

[] Yaozhi Zhang and Nina Katrine Prebensen. 2024. Co-creating with ChatGPT for tourism marketing materials. *Annals of Tourism Research Empirical Insights* 5, 1 (2024), 100124. doi:10.1016/j.annale.2024.100124

[] Zihuai Zhao, Wenqi Fan, Jiatong Li, Yunqing Liu, Xiaowei Mei, Yiqi Wang, Zhen Wen, Fei Wang, Xiangyu Zhao, Jiliang Tang, and Qing Li. 2024. Recommender Systems in the Era of Large Language Models (LLMs). *IEEE Transactions on Knowledge and Data Engineering* 36, 11 (2024), 6889–6907. doi:10.1109/TKDE.2024.3392335

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009

Temporary page!

LaTeX was unable to guess the total number of pages correctly. As there was some unprocessed data that should have been added to the final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will go away, because LaTeX now knows how many pages to expect for this document.